

Machine Learning

Lecture 3 - Evaluating predictive performance and hyperparameter learning

Mattias Villani

Department of Statistics
Stockholm University

Department of Computer and Information Science
Linköping University



Lecture overview

- **Loss functions** and evaluation metrics
- **Evaluating a classifier**
- **Estimating generalization**
- **Generalization gap**
- **Bias-variance trade-off**

Loss functions and cost functions

- Parametric model $y|x \sim f_{\theta}(y|x)$.
- Learn the parameters by minimizing a cost function

$$\hat{\theta} = \arg \min_{\theta} \underbrace{\frac{1}{n} \sum_{i=1}^n \overbrace{L(\hat{y}(x_i; \theta), y_i)}^{\text{loss function}}}_{\text{cost function } J(\theta)}$$

- **Squared error** for regression:

$$L(\hat{y}(x_i; \theta), y_i) = (\hat{y}(x_i; \theta) - y_i)^2$$

- Squared error loss = negative Gaussian log-likelihood

$$\log p(y_i|x_i, \beta) = -\frac{1}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} (y_i - x_i^{\top} \beta)^2$$

Loss functions and cost functions

- Learn the parameters by minimizing a cost function

$$\hat{\theta} = \arg \min_{\theta} \underbrace{\frac{1}{n} \sum_{i=1}^n \overbrace{L(\hat{y}(x_i; \theta), y_i)}^{\text{loss function}}}_{\text{cost function } J(\theta)}$$

- Cross-entropy** for binary classification $y \in \{0, 1\}$

$$L(\hat{y}(x_i; \theta), y_i) = -y_i \log \Pr(y_i = 1 | x_i, \theta) - (1 - y_i) \log \Pr(y_i = 0 | x_i, \theta)$$

- Cross-entropy loss = negative Bernoulli log-likelihood

$$\Pr(Y_i = y_i | x_i, \beta) = \Pr(y_i = 1 | x_i, \beta)^{y_i} \Pr(y_i = 0 | x_i, \beta)^{1-y_i}$$

where for logistic regression

$$\Pr(y_i = 1 | x_i, \beta) = \frac{\exp(x_i^T \beta)}{1 + \exp(x_i^T \beta)}.$$

- See Bonus slides for a little more on (cross-)entropy.

Detecting fraudulent banknotes

- Dataset with 1372 photographed banknotes. 610 fake.
- Raw data: $400 \times 400 = 160000$ gray-scale pixels.
- The 160000 pixel variables are condensed to four **features**:
 - ▶ variance of Wavelet Transformed image
 - ▶ skewness of Wavelet Transformed image
 - ▶ kurtosis of Wavelet Transformed image
 - ▶ entropy of image
- **Deep learning** on raw images (more later!).



Detecting fraudulent banknotes

- 1000 images for training. Predictions on 372 test images.
- Logistic regression

$$\Pr(\text{Fraud} = \text{True} | \text{features}) = \frac{\exp(\mathbf{x}^\top \beta_c)}{\sum_{j=1}^C \exp(\mathbf{x}^\top \beta_j)}$$

- **Decision**: signal fraud if $\Pr(\text{Fraud} = \text{True} | \text{Features}) > 0.5$.
- **Confusion matrix**

		Truth	
		No fraud	Fraud
Decision	No Fraud	208	1
	Fraud	3	160

Evaluating a classifier - confusion matrix

■ Confusion matrix:

		Truth	
		Positive	Negative
Decision	Positive	tp	fp
	Negative	fn	tn

- tp = true positive, fp = false positive
fn = false negative, tn = true negative.

■ Example:

		Truth	
		Fraud	No Fraud
Decision	Fraud	208	1
	No Fraud	3	160

Evaluating a classifier - Accuracy

- **Accuracy** is the proportion of correctly classified items

$$\text{Accuracy} = \frac{tp + tn}{tp + tn + fn + fp}$$

		Truth	
		Positive	Negative
Decision	Positive	tp	fp
	Negative	fn	tn

- Note: evaluation criteria (e.g. accuracy) need not be the same as the training loss function (e.g. cross-entropy).

Evaluating a classifier - Precision

- **Precision** is the proportion of truly positive items among those signaled as positive:

$$\text{Precision} = \frac{tp}{tp + fp}$$

		Truth	
		Positive	Negative
Decision	Positive	tp	fp
	Negative	fn	tn

- High precision:
 - ▶ trustworthy positives
 - ▶ people pointed out as fraudulent are almost always frauds.

Evaluating a classifier - Recall

- **Recall** is the proportion of signaled positive items among those that are truly positive:

$$\text{Recall} = \frac{tp}{tp + fn}$$

		Truth	
		Positive	Negative
Decision	Positive	tp	fp
	Negative	fn	tn

- High recall:
 - ▶ will find the positive items.
 - ▶ frauds will be caught.
- Recall is also called **True Positive Rate (TPR)** or **sensitivity**.
- There is a trade-off between Precision and Recall.

Evaluating a classifier - False Positive Rate

- **False Positive Rate (FPR)** is the proportion of signaled positive items among those that are truly negative:

$$\text{FPR} = \frac{fp}{fp + tn}$$

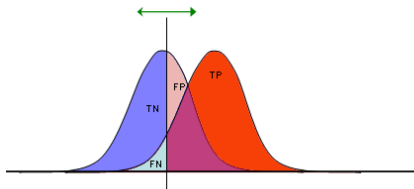
		Truth	
		Positive	Negative
Decision	Positive	tp	fp
	Negative	fn	tn

- Low FPR:
 - ▶ will very rarely signal a positive for a negative item.
 - ▶ people will not be falsely accused of fraud.


Evaluating a classifier - ROC curve

- Precision and recall depends on the **decision threshold**.
- $\Pr(\text{Spam}|\text{text in an email}) = 0.9$. Do we send it to the spam-box?
- Is $\Pr(\text{Fraud}|\text{features}) > 0.5$ a good **decision threshold**?
- **Optimal decisions** depend on the consequences.
Decision theory.
- **ROC-curve**: Receiver Operating Characteristic.
- ROC: Plots the true positive rate (TPR) against the false positive rate (FPR) **at various thresholds**.
- **AUC** = Area Under Curve. Area under the ROC curve.

Evaluating a classifier - ROC



TP	FP
FN	TN
1	1

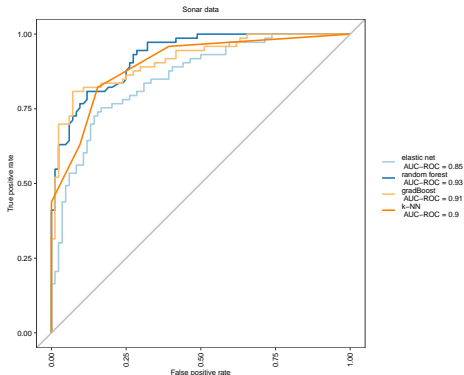


From Wikipedia.

ROC - Sonar data

■ Sonar data


- ▶ Binary response (metal/rock)
- ▶ 61 covariates (energy within a frequency band from sound impulse).



Predicting firm bankruptcy

■ Comparing test AUC for

- ▶ logistic reg (LR)
- ▶ logistic reg with additive splines
- ▶ neural networks (NN) with 2-5 layers, each with 100 nodes.



- epochs is number of iterations of the stochastic gradient decent optimization for the neural network, see Lecture 5.

Error function and generalization performance

- **Error function** compares prediction $\hat{y}(x_*)$ to actual y
 - ▶ **Squared error** for regression:

$$E(\hat{y}, y) = (\hat{y} - y)^2$$

- ▶ **Misclassification** for classification:

$$E(\hat{y}, y) = \mathbb{I}\{\hat{y} \neq y\} = \begin{cases} 0 & \text{if } \hat{y} = y \\ 1 & \text{if } \hat{y} \neq y \end{cases}$$

- Loss function $L(\hat{y}, y)$ need not be the same as error function $E(\hat{y}, y)$ used for performance evaluation.
 - ▶ k -NN does not have an explicit loss function.
 - ▶ Misclassification rate is discontinuous in θ . Hard to optimize.
- For all data: **mean squared error** and **misclassification rate**.
- **Generalization performance**: method's average performance on data drawn from a distribution $p(x, y)$.

Generalization and Training-Test splits

- How estimate generalization performance? Training-Test split.
- Partition the data in two parts:
 - ▶ **Training data** $\mathcal{T} = \{y_i, x_i\}_{i=1}^n$ to estimate parameters.
 - ▶ **Test data** (hold-out) to estimate generalization.



- How to make the split:
 - ▶ Randomly
 - ▶ Systematic (time series: most recent data in test set).
- Caret package in R. 75% training, 25% test:

```
library(caret)
library(mlbench)
data(Sonar)
inTrain <- createDataPartition(y = Sonar$Class, p = .75, list = FALSE)
training <- Sonar[ inTrain,]
testing  <- Sonar[-inTrain,]
```

Generalization and Training-Test splits

- **Training data** $\mathcal{T} = \{y_i, x_i\}_{i=1}^n$ gives estimated/trained model for prediction: $\hat{y}(x_\star; \mathcal{T})$.

- **Expected new data error**

$$E_{\text{new}} \equiv \mathbb{E}_\star [E(\hat{y}(x_\star; \mathcal{T}), y_\star)] = \int E(\hat{y}(x_\star; \mathcal{T}), y_\star) p(x_\star, y_\star) dx_\star dy_\star$$

- ▶ $E(\hat{y}(x_\star; \mathcal{T}), y_\star)$ is the error from predicting y_\star with a model trained on a **fixed training data** \mathcal{T} .
- ▶ \mathbb{E}_\star is the expectation **with respect to all possible test data** from $(x_\star, y_\star) \sim p(x, y)$.
- E_{new} measures how a trained model **generalizes** to new data from $p(x, y)$.
- Goal of ML: minimize E_{new} . But E_{new} is unknown!

Training error

- The training error is computed from $\mathcal{T} = \{y_i, x_i\}_{i=1}^n$

$$E_{\text{train}} \equiv \frac{1}{n} \sum_{i=1}^n E(\hat{y}(x_i; \mathcal{T}), y_i)$$

- Typically $E_{\text{train}} < E_{\text{new}}$, due to **overfitting** the training data.
- Error on **hold-out validation data** $\{y'_j, x'_j\}_{j=1}^{n_v}$:

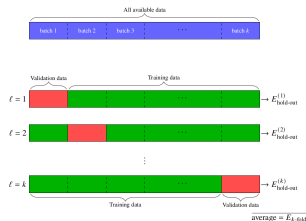
$$E_{\text{hold-out}} \equiv \frac{1}{n_v} \sum_{j=1}^{n_v} E(\hat{y}(x'_j; \mathcal{T}), y'_j)$$

- If $(y'_j, x'_j) \sim p(x, y)$, $E_{\text{hold-out}}$ is an unbiased estimate of E_{new} .
- Problem:
 - ▶ small variance of $E_{\text{hold-out}}$ requires large n_v ...
 - ▶ but large n_v means less training data ...
 - ▶ less training data means larger E_{new} . 😞

k-fold cross-validation

- Split the data in k folds or batches: B_1, \dots, B_k .
- For each batch $\ell = 1, 2, \dots, k$:
 - ▶ train the model on $B_{-\ell} \equiv \{B_1, \dots, B_{\ell-1}, B_{\ell+1}, \dots, B_k\}$.
 - ▶ predict the data in B_ℓ and compute error function $E_{\text{hold-out}}^{(\ell)}$
- Compute the estimate crossvalidation estimate of E_{new}


$$E_{\text{hold-out}} = \frac{1}{k} \sum_{\ell=1}^k E_{\text{hold-out}}^{(\ell)}$$



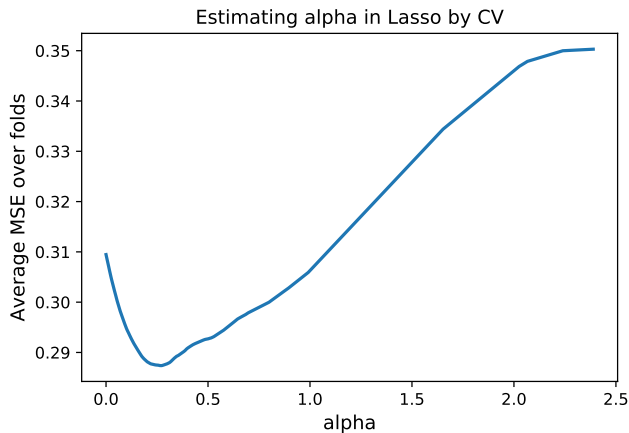
- $E_{\text{hold-out}}^{(\ell)}$ unbiased for E_{new} with $\mathcal{T} = B_{-\ell}$ for each ℓ .
- $E_{\text{hold-out}}$ (slightly) biased for E_{new} with $\mathcal{T} = \{\text{all data}\}$.

k-fold cross-validation

- **Leave-one-out** cross-validation: $k = n$. Requires n re-fits.
- Cross-validation is often used to estimate hyperparameters:
 - ▶ number of neighbors, k , in k -NN
 - ▶ tree depth in regression trees
- k -NN: $\hat{k}_{cv} = \arg \min_k E_{\text{hold-out}}(k)$.
- $E_{\text{hold-out}}(\hat{k}_{cv})$ typically underestimates E_{new} .
- Solution: set aside clean test data to estimate E_{new} .



Cross-validation learning of Lasso shrinkage



Generalization gap

- E_{new} depends on a specific \mathcal{T} . Better notation: $E_{\text{new}}(\mathcal{T})$.
- **Training-averaged generalization error:**

$$\bar{E}_{\text{new}} \equiv \mathbb{E}_{\mathcal{T}} [E_{\text{new}}(\mathcal{T})]$$

- \bar{E}_{new} is average E_{new} over all possible training data of size n .
- Training-averaged training error:

$$\bar{E}_{\text{train}} \equiv \mathbb{E}_{\mathcal{T}} [E_{\text{train}}(\mathcal{T})]$$

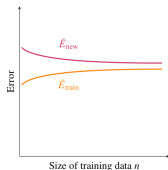
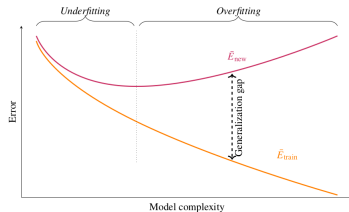
- **Generalization gap**

$$\text{generalization gap} = \bar{E}_{\text{new}} - \bar{E}_{\text{train}}$$

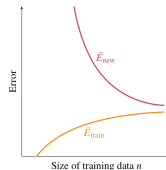
- Training error-generalization gap decomposition

$$\bar{E}_{\text{new}} = \bar{E}_{\text{train}} + \text{generalization gap}$$

Generalization gap




(a) Simple model



(b) Complex model

- $E_{\text{hold-out}} \approx E_{\text{train}}$. Possibly underfitting. Increase complexity.
- $E_{\text{train}} \approx 0$ and $E_{\text{hold-out}}$ sizeable. Possibly overfitting. Decrease complexity.

Generalization gap - five models example



Bias and Variance

- True relationship

$$y = f_0(x) + \varepsilon, \quad \mathbb{V}(\varepsilon) = \sigma^2.$$

- Estimated model $\hat{y}(x; \mathcal{T})$ (linear regression $\hat{y}(x; \mathcal{T}) = x^T \hat{\beta}$).
- Average trained model

$$\bar{f}(x) \equiv \mathbb{E}_{\mathcal{T}} [\hat{y}(x; \mathcal{T})]$$

- Bias**

$$\text{Bias}(\hat{y}(x; \mathcal{T})) = \bar{f}(x) - f_0(x)$$

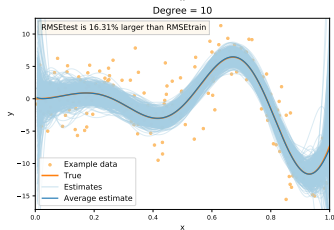
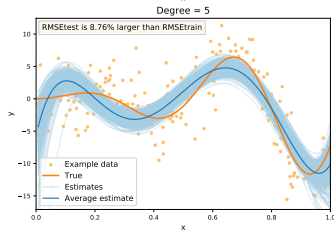
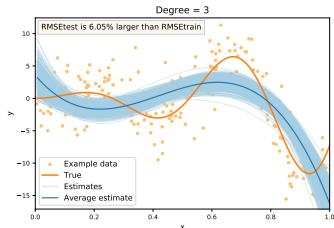
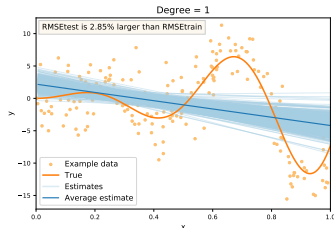
- Variance**

$$\mathbb{V}_{\mathcal{T}} [\hat{y}(x; \mathcal{T})] = \mathbb{E}_{\mathcal{T}} \left[(\hat{y}(x; \mathcal{T}) - \bar{f}(x))^2 \right].$$

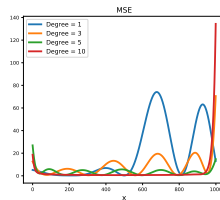
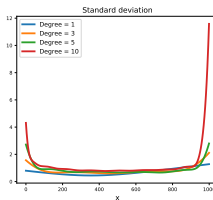
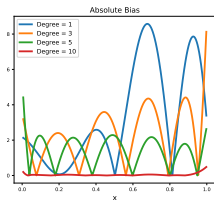
- Mean Squared Error (MSE)**

$$\text{MSE}(\hat{y}(x; \mathcal{T})) = \mathbb{E}_{\mathcal{T}} \left[(\hat{y}(x; \mathcal{T}) - f_0(x))^2 \right] = \mathbb{V}_{\mathcal{T}} [\hat{y}(x; \mathcal{T})] + \text{Bias}(\hat{y}(x; \mathcal{T}))^2$$

Bias-Variance trade-off polynomials



Bias-Variance trade-off polynomials



Bias-Variance trade-off

- Bias-variance decomposition of

$$E_{\text{new}} \equiv \mathbb{E}_{\mathcal{T}} \left[\mathbb{E}_{\star} (\hat{y}(x_{\star}; \mathcal{T}) - y_{\star})^2 \right]$$

- Change order of the expectations and insert true model:

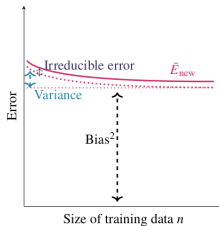
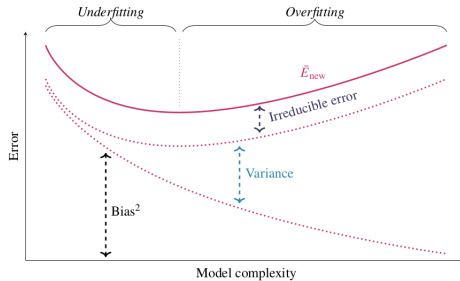
$$\bar{E}_{\text{new}} = \mathbb{E}_{\mathcal{T}} \left[\mathbb{E}_{\star} (\hat{y}(x_{\star}; \mathcal{T}) - y_{\star})^2 \right] = \mathbb{E}_{\star} \left[\mathbb{E}_{\mathcal{T}} (\hat{y}(x_{\star}; \mathcal{T}) - f_0(x_{\star}) - \varepsilon)^2 \right]$$

- Add and subtract $\bar{f}(x)$ and expand the square to get ...

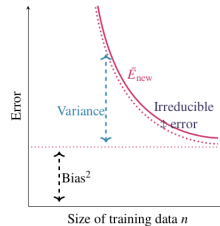
- **Bias-variance decomposition**

$$\bar{E}_{\text{new}} = \underbrace{\mathbb{E}_{\star} \left[(\bar{f}(x_{\star}) - f_0(x_{\star}))^2 \right]}_{\text{Bias}^2} + \underbrace{\mathbb{E}_{\star} \left[\mathbb{E}_{\mathcal{T}} (\hat{y}(x_{\star}; \mathcal{T}) - \bar{f}(x_{\star}))^2 \right]}_{\text{Variance}} + \underbrace{\sigma^2}_{\text{Irreducible error}}$$

Bias-Variance trade-off

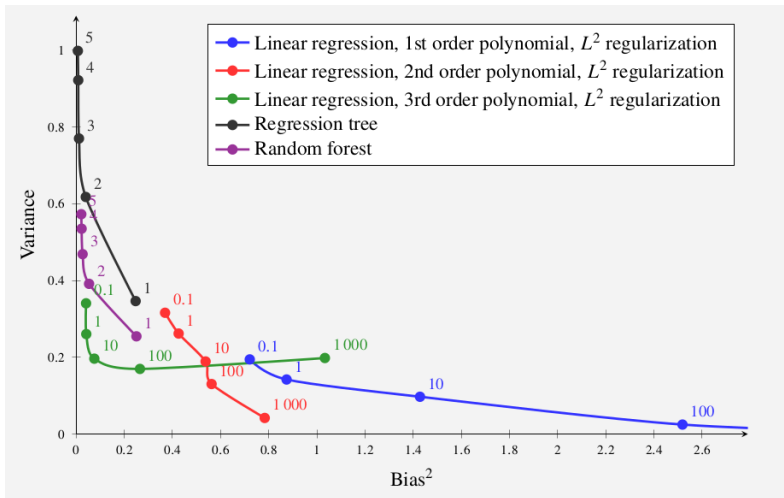


(a) Simple model

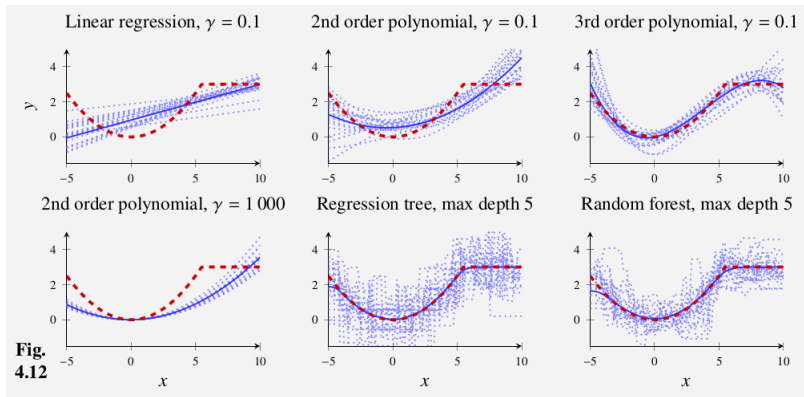


(b) Complex model

Bias-variance - five models example

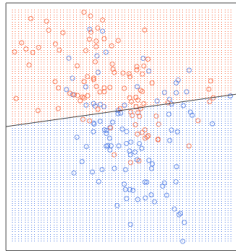


Bias-variance - five models example

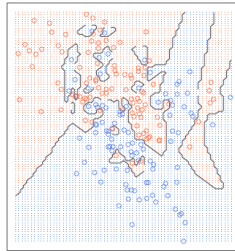


Bias-Variance trade-off classification

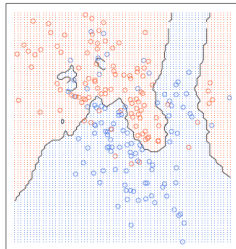
logistic regression



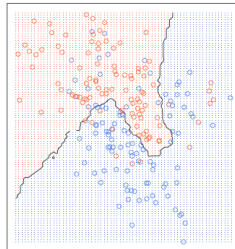
1-nearest neighbour



5-nearest neighbour

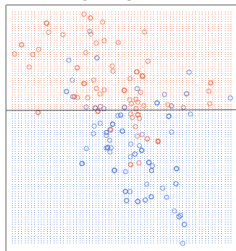


15-nearest neighbour

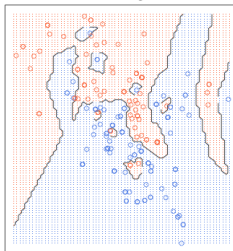


Bootstrap sample no 1

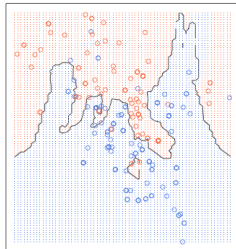
logistic regression



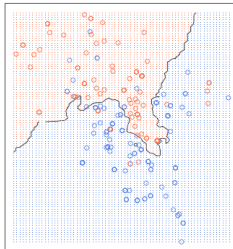
1-nearest neighbour



5-nearest neighbour

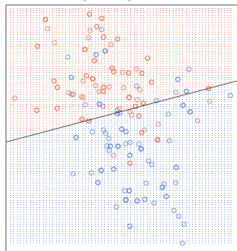


15-nearest neighbour




Bootstrap sample no 2

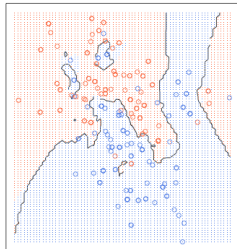
logistic regression




1-nearest neighbour



5-nearest neighbour

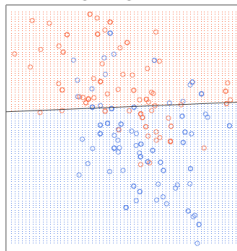


15-nearest neighbour

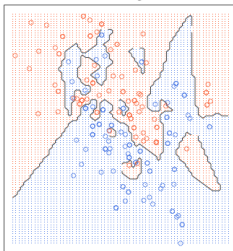


Bootstrap sample no 3

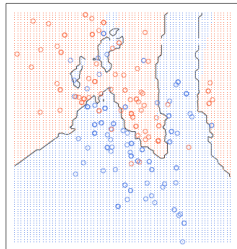
logistic regression



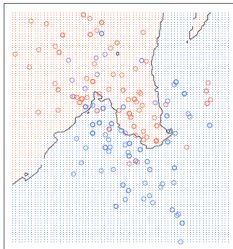
1-nearest neighbour



5-nearest neighbour

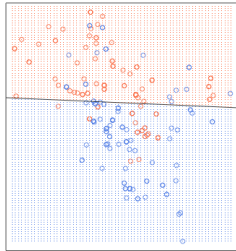


15-nearest neighbour

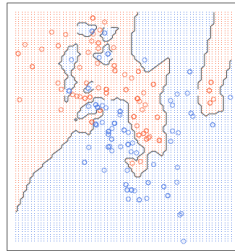


Bootstrap sample no 4

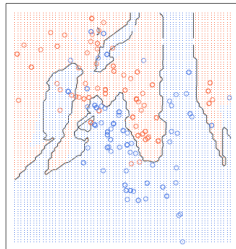
logistic regression



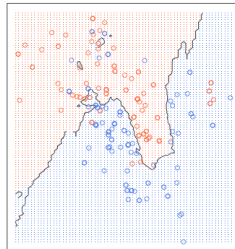
1-nearest neighbour



5-nearest neighbour

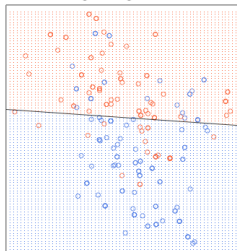


15-nearest neighbour

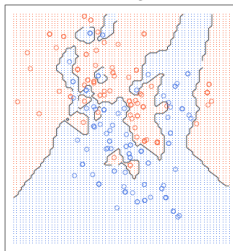


Bootstrap sample no 5

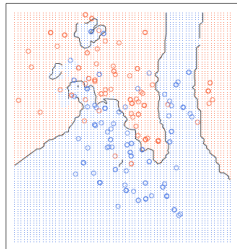
logistic regression



1-nearest neighbour



5-nearest neighbour



15-nearest neighbour

