

# Machine Learning

## Lecture 3 - Evaluating predictive performance and hyperparameter learning

**Mattias Villani**

Department of Statistics  
Stockholm University

Department of Computer and Information Science  
Linköping University



# Lecture overview

- **Loss functions** and evaluation metrics
- **Evaluating a classifier**
- **Estimating generalization**
- **Generalization gap**
- **Bias-variance trade-off**

# Loss functions and cost functions

- Parametric model  $y|x \sim f_{\theta}(y|x)$ .
- Learn the parameters by minimizing a cost function

$$\hat{\theta} = \arg \min_{\theta} \underbrace{\frac{1}{n} \sum_{i=1}^n \overbrace{L(\hat{y}(x_i; \theta), y_i)}^{\text{loss function}}}_{\text{cost function } J(\theta)}$$

- **Squared error** for regression:

$$L(\hat{y}(x_i; \theta), y_i) = (\hat{y}(x_i; \theta) - y_i)^2$$

- Squared error loss = negative Gaussian log-likelihood

$$\log p(y_i|x_i, \beta) = -\frac{1}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} (y_i - x_i^{\top} \beta)^2$$

# Loss functions and cost functions

- Learn the parameters by minimizing a cost function

$$\hat{\theta} = \arg \min_{\theta} \underbrace{\frac{1}{n} \sum_{i=1}^n \overbrace{L(\hat{y}(x_i; \theta), y_i)}^{\text{loss function}}}_{\text{cost function } J(\theta)}$$

- Cross-entropy** for binary classification  $y \in \{0, 1\}$

$$L(\hat{y}(x_i; \theta), y_i) = -y_i \log \Pr(y_i = 1 | x_i, \theta) - (1 - y_i) \log \Pr(y_i = 0 | x_i, \theta)$$

- Cross-entropy loss = negative Bernoulli log-likelihood

$$\Pr(Y_i = y_i | x_i, \beta) = \Pr(y_i = 1 | x_i, \beta)^{y_i} \Pr(y_i = 0 | x_i, \beta)^{1-y_i}$$

where for logistic regression

$$\Pr(y_i = 1 | x_i, \beta) = \frac{\exp(x_i^T \beta)}{1 + \exp(x_i^T \beta)}.$$

- See Bonus slides for a little more on (cross-)entropy.

# Evaluating a classifier - confusion matrix

## ■ Confusion matrix:

		Truth	
		Positive	Negative
Decision	Positive	tp	fp
	Negative	fn	tn

- tp = true positive, fp = false positive  
fn = false negative, tn = true negative.

## ■ Example:

		Truth	
		Fraud	No Fraud
Decision	Fraud	208	1
	No Fraud	3	160

# Evaluating a classifier - Accuracy

- **Accuracy** is the proportion of correctly classified items

$$\text{Accuracy} = \frac{tp + tn}{tp + tn + fn + fp}$$

		Truth	
		Positive	Negative
Decision	Positive	tp	fp
	Negative	fn	tn

- Note: evaluation criteria (e.g. accuracy) need not be the same as the training loss function (e.g. cross-entropy).

# Evaluating a classifier - Precision

- **Precision** is the proportion of truly positive items among those signaled as positive:

$$\text{Precision} = \frac{tp}{tp + fp}$$

		Truth	
		Positive	Negative
Decision	Positive	tp	fp
	Negative	fn	tn

- High precision:
  - ▶ trustworthy positives
  - ▶ people pointed out as fraudulent are almost always frauds.

# Evaluating a classifier - Recall

- **Recall** is the proportion of signaled positive items among those that are truly positive:

$$\text{Recall} = \frac{tp}{tp + fn}$$

		Truth	
		Positive	Negative
Decision	Positive	tp	fp
	Negative	fn	tn

- High recall:
  - ▶ will find the positive items.
  - ▶ frauds will be caught.
- Recall is also called **True Positive Rate (TPR)** or **sensitivity**.
- There is a trade-off between Precision and Recall.



# Evaluating a classifier - False Positive Rate

- **False Positive Rate (FPR)** is the proportion of signaled positive items among those that are truly negative:

$$\text{FPR} = \frac{fp}{fp + tn}$$


		Truth	
		Positive	Negative
Decision	Positive	tp	<b>fp</b>
	Negative	fn	tn

- Low FPR:
  - ▶ will very rarely signal a positive for a negative item.
  - ▶ people will not be falsely accused of fraud.


# Evaluating a classifier - ROC curve

- Precision and recall depends on the **decision threshold**.
- $\Pr(\text{Spam}|\text{text in an email}) = 0.9$ . Do we send it to the spam-box?
- Is  $\Pr(\text{Fraud}|\text{features}) > 0.5$  a good **decision threshold**?
- **Optimal decisions** depend on the consequences.  
**Decision theory**.
- **ROC-curve**: Receiver Operating Characteristic.
- ROC: Plots the true positive rate (TPR) against the false positive rate (FPR) **at various thresholds**.
- **AUC** = Area Under Curve. Area under the ROC curve.

# Evaluating a classifier - ROC



TP	FP
FN	TN
1	1




From Wikipedia.

# ROC - Sonar data

## ■ Sonar data


- ▶ Binary response (metal/rock)
- ▶ 61 covariates (energy within a frequency band from sound impulse).



# Predicting firm bankruptcy

## ■ Comparing test AUC for

- ▶ logistic reg (LR)
- ▶ logistic reg with additive splines
- ▶ neural networks (NN) with 2-5 layers, each with 100 nodes.



- epochs is number of iterations of the stochastic gradient decent optimization for the neural network, see Lecture 5.

# Error function and generalization performance

- **Error function** compares prediction  $\hat{y}(x_*)$  to actual  $y$ 
  - ▶ **Squared error** for regression:

$$E(\hat{y}, y) = (\hat{y} - y)^2$$


- ▶ **Misclassification** for classification:

$$E(\hat{y}, y) = \mathbb{I}\{\hat{y} \neq y\} = \begin{cases} 0 & \text{if } \hat{y} = y \\ 1 & \text{if } \hat{y} \neq y \end{cases}$$

- Loss function  $L(\hat{y}, y)$  need not be the same as error function  $E(\hat{y}, y)$  used for performance evaluation.
  - ▶  $k$ -NN does not have an explicit loss function.
  - ▶ Misclassification rate is discontinuous in  $\theta$ . Hard to optimize.
- For all data: **mean squared error** and **misclassification rate**.
- **Generalization performance**: method's average performance on data drawn from a distribution  $p(x, y)$ .

# Generalization and Training-Test splits

- How estimate generalization performance? Training-Test split.
- Partition the data in two parts:
  - ▶ **Training data**  $\mathcal{T} = \{y_i, x_i\}_{i=1}^n$  to estimate parameters.
  - ▶ **Test data** (hold-out) to estimate generalization.



- How to make the split:
  - ▶ Randomly
  - ▶ Systematic (time series: most recent data in test set).
- Caret package in R. 75% training, 25% test:

```
library(caret)
library(mlbench)
data(Sonar)
inTrain <- createDataPartition(y = Sonar$Class, p = .75, list = FALSE)
training <- Sonar[ inTrain,]
testing  <- Sonar[-inTrain,]
```

# Generalization and Training-Test splits

- **Training data**  $\mathcal{T} = \{y_i, x_i\}_{i=1}^n$  gives estimated/trained model for prediction:  $\hat{y}(x_\star; \mathcal{T})$ .

- **Expected new data error**

$$E_{\text{new}} \equiv \mathbb{E}_\star [E(\hat{y}(x_\star; \mathcal{T}), y_\star)] = \int E(\hat{y}(x_\star; \mathcal{T}), y_\star) p(x_\star, y_\star) dx_\star dy_\star$$

- ▶  $E(\hat{y}(x_\star; \mathcal{T}), y_\star)$  is the error from predicting  $y_\star$  with a model trained on a **fixed training data**  $\mathcal{T}$ .
- ▶  $\mathbb{E}_\star$  is the expectation **with respect to all possible test data** from  $(x_\star, y_\star) \sim p(x, y)$ .
- $E_{\text{new}}$  measures how a trained model **generalizes** to new data from  $p(x, y)$ .
- Goal of ML: minimize  $E_{\text{new}}$ . But  $E_{\text{new}}$  is unknown!



# Training error

- The training error is computed from  $\mathcal{T} = \{y_i, x_i\}_{i=1}^n$

$$E_{\text{train}} \equiv \frac{1}{n} \sum_{i=1}^n E(\hat{y}(x_i; \mathcal{T}), y_i)$$

- Typically  $E_{\text{train}} < E_{\text{new}}$ , due to **overfitting** the training data.
- Error on **hold-out validation data**  $\{y'_j, x'_j\}_{j=1}^{n_v}$ :

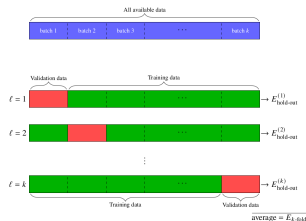
$$E_{\text{hold-out}} \equiv \frac{1}{n_v} \sum_{j=1}^{n_v} E(\hat{y}(x'_j; \mathcal{T}), y'_j)$$

- If  $(y'_j, x'_j) \sim p(x, y)$ ,  $E_{\text{hold-out}}$  is an unbiased estimate of  $E_{\text{new}}$ .
- Problem:
  - ▶ small variance of  $E_{\text{hold-out}}$  requires large  $n_v$  ...
  - ▶ but large  $n_v$  means less training data ...
  - ▶ less training data means larger  $E_{\text{new}}$ . 😞

# k-fold cross-validation

- Split the data in  $k$  folds or batches:  $B_1, \dots, B_k$ .
- For each batch  $\ell = 1, 2, \dots, k$ :
  - ▶ train the model on  $B_{-\ell} \equiv \{B_1, \dots, B_{\ell-1}, B_{\ell+1}, \dots, B_k\}$ .
  - ▶ predict the data in  $B_\ell$  and compute error function  $E_{\text{hold-out}}^{(\ell)}$
- Compute the estimate crossvalidation estimate of  $E_{\text{new}}$


$$E_{\text{hold-out}} = \frac{1}{k} \sum_{\ell=1}^k E_{\text{hold-out}}^{(\ell)}$$




- $E_{\text{hold-out}}^{(\ell)}$  unbiased for  $E_{\text{new}}$  with  $\mathcal{T} = B_{-\ell}$  for each  $\ell$ .
- $E_{\text{hold-out}}$  (slightly) biased for  $E_{\text{new}}$  with  $\mathcal{T} = \{\text{all data}\}$ .

# k-fold cross-validation

- **Leave-one-out** cross-validation:  $k = n$ . Requires  $n$  re-fits.
- Cross-validation is often used to estimate hyperparameters:
  - ▶ number of neighbors,  $k$ , in  $k$ -NN
  - ▶ tree depth in regression trees
- $k$ -NN:  $\hat{k}_{cv} = \arg \min_k E_{\text{hold-out}}(k)$ .
- $E_{\text{hold-out}}(\hat{k}_{cv})$  typically underestimates  $E_{\text{new}}$ .
- Solution: set aside clean test data to estimate  $E_{\text{new}}$ .



# Cross-validation learning of Lasso shrinkage



# Generalization gap

- $E_{\text{new}}$  depends on a specific  $\mathcal{T}$ . Better notation:  $E_{\text{new}}(\mathcal{T})$ .
- **Training-averaged generalization error:**

$$\bar{E}_{\text{new}} \equiv \mathbb{E}_{\mathcal{T}} [E_{\text{new}}(\mathcal{T})]$$

- $\bar{E}_{\text{new}}$  is average  $E_{\text{new}}$  over all possible training data of size  $n$ .
- Training-averaged training error:

$$\bar{E}_{\text{train}} \equiv \mathbb{E}_{\mathcal{T}} [E_{\text{train}}(\mathcal{T})]$$



- **Generalization gap**

$$\text{generalization gap} = \bar{E}_{\text{new}} - \bar{E}_{\text{train}}$$


- Training error-generalization gap decomposition

$$\bar{E}_{\text{new}} = \bar{E}_{\text{train}} + \text{generalization gap}$$

# Generalization gap



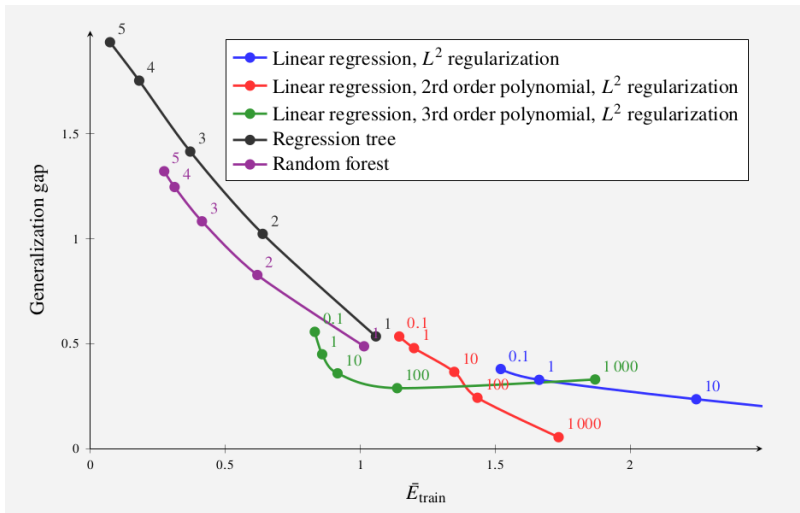
(a) Simple model



(b) Complex model

- $E_{\text{hold-out}} \approx E_{\text{train}}$ . Possibly underfitting. Increase complexity.
- $E_{\text{train}} \approx 0$  and  $E_{\text{hold-out}}$  sizeable. Possibly overfitting. Decrease complexity.

# Generalization gap - five models example



# Bias and Variance

- True relationship

$$y = f_0(x) + \varepsilon, \quad \mathbb{V}(\varepsilon) = \sigma^2.$$

- Estimated model  $\hat{y}(x; \mathcal{T})$  (linear regression  $\hat{y}(x; \mathcal{T}) = x^T \hat{\beta}$ ).
- Average trained model

$$\bar{f}(x) \equiv \mathbb{E}_{\mathcal{T}} [\hat{y}(x; \mathcal{T})]$$

- Bias**

$$\text{Bias}(\hat{y}(x; \mathcal{T})) = \bar{f}(x) - f_0(x)$$

- Variance**





$$\mathbb{V}_{\mathcal{T}} [\hat{y}(x; \mathcal{T})] = \mathbb{E}_{\mathcal{T}} \left[ (\hat{y}(x; \mathcal{T}) - \bar{f}(x))^2 \right].$$

- Mean Squared Error (MSE)**




$$\text{MSE}(\hat{y}(x; \mathcal{T})) = \mathbb{E}_{\mathcal{T}} \left[ (\hat{y}(x; \mathcal{T}) - f_0(x))^2 \right] = \mathbb{V}_{\mathcal{T}} [\hat{y}(x; \mathcal{T})] + \text{Bias}(\hat{y}(x; \mathcal{T}))^2$$



# Bias-Variance trade-off polynomials



# Bias-Variance trade-off polynomials



# Bias-Variance trade-off

- Bias-variance decomposition of

$$E_{\text{new}} \equiv \mathbb{E}_{\mathcal{T}} \left[ \mathbb{E}_{\star} (\hat{y}(x_{\star}; \mathcal{T}) - y_{\star})^2 \right]$$

- Change order of the expectations and insert true model:



$$\bar{E}_{\text{new}} = \mathbb{E}_{\mathcal{T}} \left[ \mathbb{E}_{\star} (\hat{y}(x_{\star}; \mathcal{T}) - y_{\star})^2 \right] = \mathbb{E}_{\star} \left[ \mathbb{E}_{\mathcal{T}} (\hat{y}(x_{\star}; \mathcal{T}) - f_0(x_{\star}) - \varepsilon)^2 \right]$$

- Add and subtract  $\bar{f}(x)$  and expand the square to get ...


- **Bias-variance decomposition**

$$\bar{E}_{\text{new}} = \underbrace{\mathbb{E}_{\star} \left[ (\bar{f}(x_{\star}) - f_0(x_{\star}))^2 \right]}_{\text{Bias}^2} + \underbrace{\mathbb{E}_{\star} \left[ \mathbb{E}_{\mathcal{T}} (\hat{y}(x_{\star}; \mathcal{T}) - \bar{f}(x_{\star}))^2 \right]}_{\text{Variance}} + \underbrace{\sigma^2}_{\text{Irreducible error}}$$

# Bias-Variance trade-off




(a) Simple model




(b) Complex model

# Bias-variance - five models example




# Bias-variance - five models example




# Bias-Variance trade-off classification


**logistic regression**




**1-nearest neighbour**



**5-nearest neighbour**




**15-nearest neighbour**




# Bootstrap sample no 1


logistic regression




1-nearest neighbour



5-nearest neighbour




15-nearest neighbour






# Bootstrap sample no 2


**logistic regression**




**1-nearest neighbour**



**5-nearest neighbour**




**15-nearest neighbour**




# Bootstrap sample no 3


**logistic regression**




**1-nearest neighbour**



**5-nearest neighbour**




**15-nearest neighbour**




# Bootstrap sample no 4


**logistic regression**




**1-nearest neighbour**



**5-nearest neighbour**




**15-nearest neighbour**




# Bootstrap sample no 5


**logistic regression**



**1-nearest neighbour**



**5-nearest neighbour**



**15-nearest neighbour**

