

# Machine Learning

## Lecture 10 - Textual data, document classification and topic models

**Mattias Villani**

Department of Statistics  
Stockholm University

Department of Computer and Information Science  
Linköping University



# Lecture overview

- Text as data
- Classifying documents
- Topic models

# Text is data

- **Digitalization**: text is becoming an important data source.
- The web, PDF documents (legal, political, medical, etc)
- **Unstructured** (not tables), yet **structured** (by language).
- **Big data**. 100K, 1M, 1B documents in a data set.
- **Pre-processing** to get data useful for statistical analysis.
- **Feature construction** - turning text into numbers.

# Text applications

- **Language models** (predict the next word on smartphone)
- **Machine translation** (Google translate)
- **Document classification** (Shakespeare? Spam and blog filters. harmful EULA)
- **Sentiment analysis** (positive/negative sentiment in tweets or financial statements)
- **Information retrieval** (Google search)
- **Part-of-speech tagging** (predict grammatical category)
- **Prediction models** based on text.
  - Predicting financial turbulence from economic press.
  - Finding bugs from bug reports

# Classifying texts - feature construction

- Data: **corpus** of **documents**, each with a **label**
  - ▶ Journal articles under the headings: sports, politics, culture etc
  - ▶ Financial articles: positive/negative about the economy.
- **Text features** from documents:
  - ▶ Presence/absence of individual words, or pairs of words
  - ▶ Number of times an individual word is used
  - ▶ Lexical diversity
  - ▶ Number of web links from document (Page Rank).

| Document | has('ball') | has('political_arena') | wordlen | Lex. Div. | Topic  |
|----------|-------------|------------------------|---------|-----------|--------|
| Article1 | Yes         | No                     | 4.1     | 5.4       | Sports |
| Article2 | No          | No                     | 6.5     | 13.4      | Sports |
| ⋮        | ⋮           |                        | ⋮       | ⋮         | ⋮      |
| ArticleN | No          | Yes                    | 7.4     | 11.1      | News   |

- See `DocumentTermMatrix` function in `tm` package.

# Multinomial model with Dirichlet prior

- **Topic model.** Factorization of the **Document-Term matrix**.
- **Categorical counts:**  $y = (y_1, \dots, y_C)$ , where  $\sum_{c=1}^C y_c = n$ .
- $y_c$  = number of observations in  $c$ th category. Brand choices.
- **Multinomial model:**

$$p(y|\theta) \propto \prod_{c=1}^C \theta_c^{y_c}, \text{ where } \sum_{c=1}^C \theta_c = 1.$$

- **Dirichlet prior:**  $\theta \sim \text{Dirichlet}(\alpha_1, \dots, \alpha_C)$


$$p(\theta) \propto \prod_{c=1}^C \theta_c^{\alpha_c - 1}.$$

# Dirichlet prior

$$(\theta_1, \dots, \theta_C) \sim \text{Dirichlet}(\alpha_1, \dots, \alpha_C)$$

$$\mathbb{E}(\theta_c) = \frac{\alpha_c}{\sum_{j=1}^C \alpha_j}$$

$$\mathbb{V}(\theta_c) = \frac{\mathbb{E}(\theta_c)(1 - \mathbb{E}(\theta_c))}{1 + \sum_{j=1}^C \alpha_j}$$



- Uniform distribution on unit simplex:  $\alpha_1 = \dots = \alpha_K = 1$ .

# Multinomial model with Dirichlet prior

## Multinomial data with Dirichlet prior

**Model:**  $\mathbf{n}|\boldsymbol{\theta} \sim \text{Multinomial}(\boldsymbol{\theta})$ , where  
 $\mathbf{n} = (n_1, \dots, n_C)$  are counts in  $C$  categories  
 $\boldsymbol{\theta} = (\theta_1, \dots, \theta_C)$  are category probabilities.

**Prior:**  $\boldsymbol{\theta} \sim \text{Dirichlet}(\boldsymbol{\alpha})$ , for  $\boldsymbol{\alpha} = (\alpha_1, \dots, \alpha_C)$





**Posterior:**  $\boldsymbol{\theta} \sim \text{Dirichlet}(\boldsymbol{\alpha} + \mathbf{n})$



# Example: smartphone market shares

- Survey among 513 smartphones owners:
  - ▶ 180 used mainly an iPhone
  - ▶ 230 used mainly an Android phone
  - ▶ 62 used mainly a Windows phone
  - ▶ 41 used mainly some other mobile phone.
- Old survey: iPhone 30%, Android 30%, Windows 20%, Other 20%.
- **Pr(Android has largest share | Data)**
- Prior:  $\alpha_1 = 15, \alpha_2 = 15, \alpha_3 = 10$  and  $\alpha_4 = 10$  (prior info is equivalent to a survey with only 50 respondents)
- Posterior:  $(\theta_1, \theta_2, \theta_3, \theta_4) | y \sim \text{Dirichlet}(195, 245, 72, 51)$ .
- **R Notebook:** [Multinomial.Rmd](#)

# Example: smartphone market shares




# Mixture of unigrams

- Let  $\phi_1, \phi_2, \dots, \phi_K$  be distributions over the vocabulary. **Topics**.

| Topic | Word distr. | probability | dna | gene | data | distribution |
|-------|-------------|-------------|-----|------|------|--------------|
| 1     | $\phi_1$    | 0.5         | 0.1 | 0.0  | 0.2  | 0.2          |
| 2     | $\phi_2$    | 0.0         | 0.5 | 0.4  | 0.1  | 0.0          |

- For each document  $d = 1, \dots, D$ :
  - Draw a **topic**  $z_d$  from a **topic distribution**  $\theta = (\theta_1, \dots, \theta_K)$ .
  - Given topic  $z_d$ , draw **words** from a **word distribution**  $\phi_{z_d}$ .



- Each document belongs to **exactly** one topic.
- Topic models are **mixed-membership models**.

# Simulating documents from a topic model

■ Assume that we have:

- ▶ A fixed vocabulary  $V$
- ▶  $D$  documents
- ▶  $N$  words in each document
- ▶  $K$  topics

1 For each **topic** ( $k = 1, \dots, K$ ):

- a. Draw a distribution over the words  $\phi_k \sim \text{Dir}(\eta, \eta, \dots, \eta)$

2 For each **document** ( $d = 1, \dots, D$ ):

- a. Draw a vector of topic proportions  $\theta_d \sim \text{Dir}(\alpha_1, \dots, \alpha_K)$
- b. For each **word** ( $i = 1, \dots, N$ ):
  - i. Draw a topic assignment  $z_{di} \sim \text{Categorical}(\theta_d)$
  - ii. Draw a word  $w_{di} \sim \text{Categorical}(\phi_{z_{di}})$

# Example - simulation from two topics

| Topic | Word distr. | probability | dna | gene | data | distribution |
|-------|-------------|-------------|-----|------|------|--------------|
| 1     | $\phi_1$    | 0.5         | 0.1 | 0.0  | 0.2  | 0.2          |
| 2     | $\phi_2$    | 0.0         | 0.5 | 0.4  | 0.1  | 0.0          |

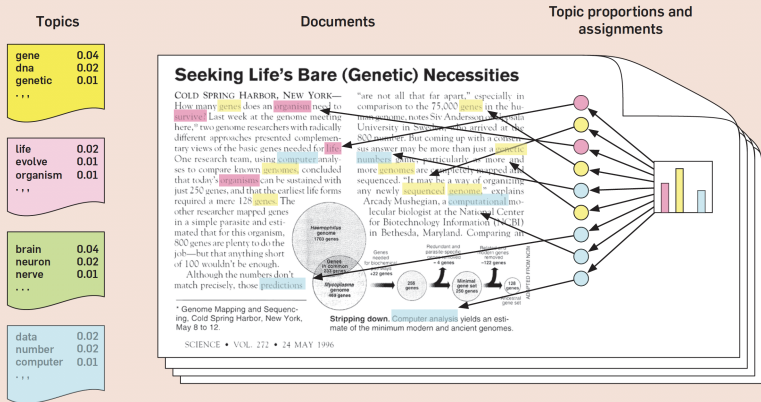
|       |                         |         |             |
|-------|-------------------------|---------|-------------|
| Doc 1 | $\theta_1 = (0.2, 0.8)$ |         |             |
|       | Word 1:                 | Topic=2 | Word='gene' |
|       | Word 2:                 | Topic=2 | Word='gene' |
|       | Word 3:                 | Topic=1 | Word='data' |

|       |                         |         |                    |
|-------|-------------------------|---------|--------------------|
| Doc 2 | $\theta_2 = (0.9, 0.1)$ |         |                    |
|       | Word 1:                 | Topic=1 | Word='probability' |
|       | Word 2:                 | Topic=1 | Word='data'        |
|       | Word 3:                 | Topic=1 | Word='probability' |

|       |                         |   |   |
|-------|-------------------------|---|---|
| Doc 3 | $\theta_2 = (0.5, 0.5)$ |   |   |
| ⋮     | ⋮                       | ⋮ | ⋮ |

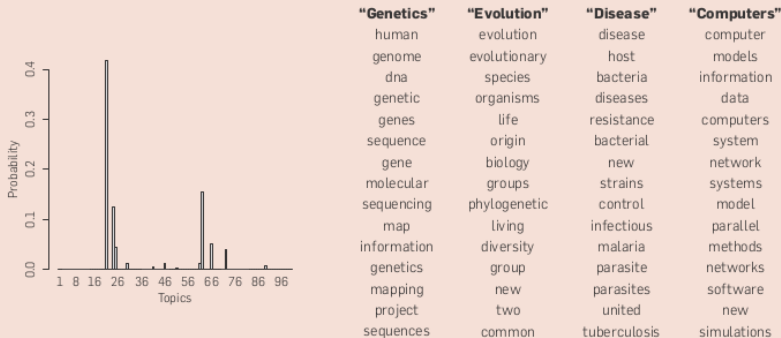
# Example from Science (Blei, review paper)

Figure 1. The intuitions behind latent Dirichlet allocation. We assume that some number of “topics,” which are distributions over words, exist for the whole collection (far left). Each document is assumed to be generated as follows. First choose a distribution over the topics (the histogram at right); then, for each word, choose a topic assignment (the colored coins) and choose the word from the corresponding topic. The topics and topic assignments in this figure are illustrative—they are not fit from real data. See Figure 2 for topics fit from data.



# Example from Science (Blei, review paper)

Figure 2. Real inference with LDA. We fit a 100-topic LDA model to 17,000 articles from the journal *Science*. At left are the inferred topic proportions for the example article in Figure 1. At right are the top 15 most frequent words from the most frequent topics found in this article.



# Predicting bug location from bug reports

```
def Network(networkInputs):
    # CODE
    # MORE CODE
    # TOO MUCH CODE

    return(networkOutputs)

#

def UI(UIInputs):
    # CODE
    # MORE CODE
    # TOO MUCH CODE


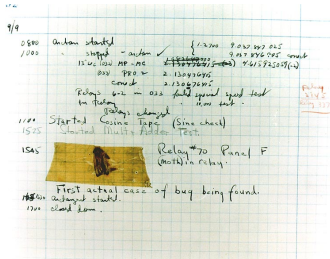
    return(UIOutputs)

#

def Database(DBInputs):
    # CODE
    # MORE CODE
    # TOO MUCH CODE

    return(DBOutputs)

#
```





# Three datasets

| Dataset        | No. Bug reports | No. classes | Vocabulary size |
|----------------|-----------------|-------------|-----------------|
| Mozilla        | 15,000          | 118         | 3505            |
| Eclipse        | 15,000          | 49          | 3367            |
| <b>Telecom</b> | <b>9,778</b>    | <b>26</b>   | <b>5286</b>     |


---

From Jonsson et al (2016). Automatic Localization of Bugs to Faulty Components in Large Scale Software Systems Using Bayesian Classification.

- Automatically summarize a bug report by topics.

| Topic | Topic label        | Top 10 words in topic   |
|-------|--------------------|---|
| 11    | HTTP               | proxy server http network connection request connect error www host |
| 27    | Layout             | div style px background color border css width height element html  |
| 28    | Connection Headers | http cache accept en public localhost gmt max modified alive        |
| 55    | Search             | search google bar results box type find engine enter text           |
| 82    | Scrolling          | scroll page scrolling mouse scrollbar bar left bottom click content |

# Topic proportions $\theta_d$



# Multi-class logistic regression on topics

